



RTL And SystemVerilog

📅 Date	07/05/2026
📌 Note Type	Notes

Key Takeaways

1. RTL describes how data flows between **registers** (state elements) through **combinational logic** (transfers) on each clock cycl.
2. Serves as the "golden contract" between **Design Verification (DV)** and **Physical Design (PD)** teams.
3. **Design Abstraction Levels**
4. **RTL-to-Silicon Flow (RTL → Synthesis → Place & Route → GDSII)**
5. **Concurrent vs. Procedural Code**
6. **Operator Types**
7. RTL design with SystemVerilog is about **describing hardware structure**, not writing software. It requires understanding concurrency, cycle-accurate behavior, and the distinction between code that can become real silicon versus code that only helps simulate and verify.
8. Concept learning for **Sequence Detector**, **Traffic Light controller**, and **FIFO design**
9. Verification planning for DV.

Action Items

- ☐ Introduction to RTL design
- ☐ Learn about SystemVerilog
- ☐ Covering fundamentals of hardware description
- ☐ SV coding concept
- ☐ Sequence Detector
- ☐ Traffic Light Controller
- ☐ FIFO Design
- ☐ Verification Planning

Notes

Understanding RTL and SV Essentials

1.1 What is RTL?

Register Transfer Level (RTL) is the fundamental abstraction layer where modern digital design and verification happen. It bridges the gap between a high-level algorithmic concept and the physical transistors on a silicon wafer.

Definition of Register Transfer Level (RTL)

The term "**register transfer**" describes exactly how data flows through a synchronous digital circuit.

In RTL hardware design, you are defining two main things:

1. **Registers (State Elements):** Typically flip-flops that hold data and update their values on a clock edge.
2. **Transfers (Combinational Logic):** The Boolean logic gates (adders, multiplexers, multipliers) that manipulate the data as it travels from the output of one register to the input of the next.

Instead of drawing individual logic gates, you write code that dictates how data is routed and transformed between these memory elements at every clock cycle.

Abstraction Levels: Behavioral vs. RTL vs. Gate-Level

To understand hardware design, it helps to see where RTL sits within the design hierarchy:

Abstraction level	What it focuses on	Key characteristics / notes	Concrete example
Behavioral level	Algorithms and <i>what</i> the circuit should do	Highest abstraction; often ignores clock cycles and hardware constraints; resembles software (C/C++/Python); used for early architectural modeling.	C-like model: <code>y = (a + b) * c;</code> (no clock, no registers)
RTL (Register Transfer Level)	Cycle-accurate hardware behavior	• Middle ground; • hardware-aware; • describes what happens on each clock tick; • explicitly defines registers and combinational logic between them	SystemVerilog: <code>always_ff</code> <code>@(posedge clk) q <= d;</code> • <code>assign d = a & b;</code>
Gate-level netlist	Actual gates and flip-flops in a specific technology	Lowest level; tool-generated netlist of gates (AND/OR/NOT/DFF) mapped to a standard-cell library; shows how RTL becomes physical hardware blocks	Netlist style: <code>DFF_X1 u1 (.D(d), .CK(clk), .Q(q));</code> and <code>AND2_X1 u2 (.A(a), .B(b), .Z(d));</code>

The Standard Handoff Format

RTL serves as the "golden contract" bridging different engineering disciplines in the semiconductor industry.

- **For Design Verification (DV):** RTL provides the exact functional intent of the chip. It is abstract enough that one can write complex SystemVerilog `testbenches`, `assertions`, and `coverage metrics` to `verify its logic` without worrying about physical wiring.
- **For Physical Design (PD):** RTL is structured enough that automated tools can predictably map it into physical gates. It provides the definitive starting point for floorplanning, synthesis, and routing.

Because both teams rely on it—one for functional correctness and the other for physical implementation—RTL is the standard handoff format.

The RTL-to-Silicon Flow

The journey from SystemVerilog code to a manufactured chip follows a strict physical design pipeline:

1. **RTL:** The human-readable hardware description is finalized and verified.
2. **Synthesis:** A synthesis tool maps the RTL logic into a gate-level netlist using standard cells provided by a foundry, optimizing for initial Power, Performance, and Area (PPA) goals.
3. **Place & Route (P&R):** The synthesized gates are physically arranged on the chip's floorplan (Placement) and wired together (Routing) while strictly adhering to timing paths and Design Rule Checks (DRC).
4. **GDSII:** Once physical verification and sign-off (including static timing analysis) are complete, the final physical layout is exported as a GDSII file. This is the binary database sent to the foundry to manufacture the photomasks for fabrication.

Simulator vs. Synthesizer

When we write SystemVerilog, two completely different types of software will read our code, each with its own purpose:

- **A Simulator (used in DV):** Reads RTL to simulate *time and behavior*.
 - It executes the code sequentially to mimic how the hardware will act cycle-by-cycle.
 - Its goal is functional verification—checking if the logic produces the correct output for a given input sequence.
 - It does not care if the code is physically possible to build.
- **A Synthesizer (used in PD):** Reads RTL to infer *actual hardware*.
 - It translates the code into physical logic gates, multiplexers, and flip-flops.
 - Its goal is realization—ensuring the design can actually be manufactured and will meet strict timing and area constraints.
 - Synthesizers will reject or ignore code that cannot be mapped to hardware (such as initial blocks or infinite loops).

1.2 HDL fundamentals

What is an HDL (and why it isn't "programming")?

A **Hardware Description Language (HDL)** doesn't describe a sequence of instructions for a CPU to run. Instead, it describes **hardware structure and behavior**—signals, registers, and the logic that connects them.

- **Programming languages (C, Python):** express a *sequence* of operations executed over time on a fixed processor.
- **HDLs (Verilog, SystemVerilog):** describe *concurrent* hardware blocks that exist at the same time in space. Writing HDL is closer to drawing a schematic—just with text.

**** Concurrent:** many things happen at the same time. In **software**, lines usually run **one after another** (step 1, then step 2). In **hardware**, different parts are like **separate little machines** wired together, so they all "work" **in parallel**. Examples: If we build **five adders** in hardware, all five add at once, because they're separate circuits

Verilog vs. SystemVerilog

- **Verilog (1984):** started as a simulation language and later became widely used for synthesis; C-like syntax but limited verification capabilities.
- **SystemVerilog (2005):** a backwards-compatible superset that unified **design** and **verification**.
 - **Design improvements:** `logic`, `always_comb`, `always_ff`, better data types (structs, enums), arrays, interfaces.
 - **Verification features:** classes (OOP), constrained-random stimulus, functional coverage, and **SystemVerilog Assertions (SVA)**.

Synthesizable vs. non-synthesizable code

Because SystemVerilog is used for both RTL design and testbenches, it's helpful to separate code into:

1. **Synthesizable subset:** constructs that tools can map into real gates/flip-flops (e.g., `always_ff`, `assign`, arithmetic/boolean operators).
2. **Non-synthesizable subset:** simulator-only constructs used in testbenches (e.g., `#10` delays, `initial` blocks, `$display`, `class`-based TB code).

Delta cycles and the event queue

Simulators approximate hardware concurrency using an **event queue** with micro-steps called **delta cycles (Δ)**.


- A delta cycle advances *zero* simulation time (e.g., still at 0 ns).
- It lets the simulator resolve cause-and-effect ordering without moving the clock forward—for example, when a change at Gate A schedules an update at Gate B for the next delta cycle (`T +  $\Delta$`).

Delta cycles (Δ) and the event queue are how a simulator pretends lots of hardware things happen "at the same time," even though the simulator is running step-by-step on a computer.

- The **event queue** is like a **to-do list** the simulator uses: "Update this wire," "Recalculate that gate," etc.
- A **delta cycle (Δ)** is a **tiny step with zero time passing** (time stays at, say, 0 ns).
- Why do this? Because in real hardware, when one signal changes, it can immediately cause other signals to change. The simulator needs a clean way to handle those "instant" domino effects without moving the clock forward.

Simple picture:

When Signal A changes, the simulator schedules Signal B's update for **the next delta cycle** (written like `T +  $\Delta$`). Then it keeps doing these zero-time mini-steps until everything settles, and only *then* it moves on in real time (like to the next nanosecond or clock edge).



1.3 SystemVerilog Module Structure

Module / Endmodule

The `module` is the fundamental building block of hardware design, acting like a black box with defined pins (ports).

Code snippet

```

module my_gate (
  // Port declarations go here
);
  // Internal logic goes here
endmodule

```

Port Declarations

Ports define how the module interacts with the outside world:

- `input`: Data flows strictly *into* the module.
- `output`: Data flows strictly *out of* the module.
- `inout`: Bi-directional port (used typically for tristate pads/buses, rare inside internal RTL logic).

Data Types: Logic, Wire, Reg

Legacy Verilog required you to understand whether a signal was a structural connection (`wire`) or handled in a procedural block (`reg`). SystemVerilog simplifies this significantly.

Data Type	Description / Usage
<code>wire</code>	Verilog Legacy. Represents a physical wire. Cannot hold a state; must be continuously driven by an <code>assign</code> statement or a gate output.
<code>reg</code>	Verilog Legacy. Represents a variable used inside a procedural block (<code>always</code>). <i>Misleading name:</i> It does not necessarily infer a physical register.
<code>logic</code>	SystemVerilog Standard. Replaces both <code>wire</code> and <code>reg</code> in 99% of design code. The compiler automatically determines if it represents a wire or a register based on how it is assigned. (Note: <code>logic</code> is a 4-state type: <code>0</code> , <code>1</code> , <code>X</code> (unknown), <code>Z</code> (high-impedance)).

Parameter vs. Localparam

Used for making code highly reusable and configurable (e.g., changing data bus widths without rewriting code).

- `parameter`: Can be overridden when the module is instantiated by a higher-level module (e.g., setting a FIFO depth to 16 in one place and 64 in another).
- `localparam`: A local constant. It **cannot** be modified outside the module. Used for internal constants like State Machine encodings.

Multi-bit Signals and Part-Selects

Signals can be grouped into vectors (buses). The syntax is standard `[MSB:LSB]` (Most Significant Bit to Least Significant Bit).

Code snippet

```
logic [7:0] data_bus;      // An 8-bit wide bus (bits 7 down to 0)
logic [2:0] sub_slice;

assign sub_slice = data_bus[4:2]; // Part-select: extracts bits 4, 3, and 2
```

Signed vs. Unsigned Signals

By default, all hardware signals are **unsigned**.

- `logic [7:0] a;` → Ranges from 0 to 255.
- `logic signed [7:0] b;` → Uses 2's complement notation. Ranges from -128 to +127.
- *Warning:* If you mix signed and unsigned signals in an expression, the entire expression automatically converts to **unsigned**, which can cause critical bugs.

1.4 Concurrent vs. Procedural Code

Assign Statements

Used for **Continuous Assignment**. It models pure combinational logic outside of procedural blocks. The left-hand side is updated immediately whenever any signal on the right-hand side changes.

Code snippet

```
assign out = a & b; // Acts exactly like an AND gate wired to 'a' and 'b'
```

always_comb

A SystemVerilog procedural block dedicated exclusively to **combinational logic**.

- It re-evaluates any variable inside the block whenever its value changes.
- It guarantees that simulators will flag errors if you accidentally create a hardware latch.

always_ff

A SystemVerilog procedural block dedicated exclusively to modeling **sequential logic** (Flip-Flops/Registers).

- It requires a specific clock edge or reset edge trigger.

Code snippet

```
always_ff @(posedge clk or negedge rst_n) begin
    if (!rst_n)
        q <= 0; // Sequential assignments use non-blocking '<='
    else
        q <= d;
end
```

always_latch

Used to model level-sensitive latches. In modern synchronous digital design, **latches are generally unintended** and dangerous. They usually occur because a designer forgot to define an `else` condition in an `if-else` statement or a

`default` case in a `case` statement inside an `always_comb` block, causing the hardware to "remember" its old value.

Sensitivity Lists

The condition inside the `@(...)` parenthetical statement that tells the simulator when to execute the procedural block.

- `@(posedge clk)` : Executes only when the clock transitions from 0 to 1.
- `@(*)` or `always_comb` : Automatically includes every input signal in the block to ensure it acts like pure combinational logic.

1.5 Operators & Expressions

Assume `logic [3:0] a = 4'b1010;` and `logic [3:0] b = 4'b0110;`

Bitwise Operators

Performs the operation on each individual bit independently.

- Bitwise AND (`&`): `a & b` $\rightarrow$ `4'b0010`
- Bitwise OR (`|`): `a | b` $\rightarrow$ `4'b1110`
- Bitwise XOR (`^`): `a ^ b` $\rightarrow$ `4'b1100`
- Bitwise NOT (`~`): `~a` $\rightarrow$ `4'b0101`

Logical Operators

Evaluates expressions to a single-bit boolean truth value (`1'b1` for True, `1'b0` for False).

- Logical AND (`&&`): Returns True if both sides are non-zero.
- Logical OR (`||`): Returns True if either side is non-zero.
- Logical NOT (`!`): Inverts truth value. `!a` $\rightarrow$ `1'b0` (since `a` is non-zero).

Reduction Operators

Performs a bitwise operation across all the bits of a *single* vector, collapsing it into a 1-bit result.

- Reduction AND (`&a`): `1 & 0 & 1 & 0` $\rightarrow$ `1'b0`
- Reduction OR (`|a`): `1 | 0 | 1 | 0` $\rightarrow$ `1'b1`
- Reduction XOR (`^a`): Useful for checking parity.

Relational & Equality Operators

- `==` and `!=` : Logical equality. If any bit is `X` or `Z`, the result will return `X` (unknown).
- `===` and `!==` : Case equality (used in verification, **non-synthesizable**). Matches literal `X` and `Z` states exactly.

Concatenation & Replication

- **Concatenation** `{}` : Joins smaller signals together to form a wider bus. **Code snippet**

```
logic [1:0] x = 2'b11;
logic [1:0] y = 2'b00;
logic [3:0] z = {x, y}; // z becomes 4'b1100
```

- **Replication** `{N{}}` : Duplicates a signal `N` times. **Code snippet**

```
logic [3:0] pad = {4{1'b1}}; // pad becomes 4'b1111
```

Shift Operators

- **Logical Shifts** (`<<`, `>>`) : Shifts bits and pads empty spots with zeros.
- **Arithmetic Shifts** (`<<=`, `>>=`) : Shifts bits but preserves the **sign bit** (the MSB) when shifting right, maintaining proper 2's complement value (only works if the variable is explicitly declared as `signed`).

Here are a few clear **SystemVerilog shift-operator code examples** (logical and arithmetic), using both constant and variable shift amounts.

1) Logical shifts (<< , >>) — shift in zeros

```
module shift_logical;
  logic [7:0] a;
  logic [7:0] y1, y2;

  initial begin
    a = 8'b1011_0101; // 0xB5

    y1 = a << 2;        // 1101_0100 (0xD4) left shift, zeros on right
    y2 = a >> 3;        // 0001_0110 (0x16) right shift, zeros on left
  end
endmodule
```

2) Arithmetic right shift (>>>) — preserves sign bit (requires signed)

```
module shift_arithmetic;
  logic signed [7:0] s;
  logic signed [7:0] y;

  initial begin
    s = -8'sd8;          // 1111_1000
    y = s >>> 2;          // 1111_1110 (still negative; sign-extended)
  end
endmodule
```

3) Variable shift amount (useful in RTL)

```
module shift_variable;
  logic [15:0] data;
  logic [3:0] shamt;
  logic [15:0] out;

  always_comb begin
    out = data >> shamt; // logical right shift by shamt (0..15)
  end
endmodule
```

What is a Finite State Machine

A **Finite State Machine (FSM)** is a mathematical model of computation used to design sequential logic circuits. It describes a system that can exist in exactly **one of a finite number of states** at any given time. The machine changes from one state to another (a transition) in response to external inputs and/or clock edges.

The Four Core Components of an FSM

Every FSM, regardless of complexity, consists of exactly four elements:

1. **States:** The distinct conditions or modes the system can reside in (e.g., **IDLE**, **READ**, **WRITE**).
2. **Inputs:** External signals that trigger state changes or affect outputs.
3. **Outputs:** Signals generated by the FSM to control other parts of the hardware.
4. **Transitions:** The rules/logic defining how the FSM moves from a current state to a next state.

Why FSMs Dominate Digital Design Control Logic

While data paths handle data manipulation (like arithmetic units), FSMs act as the **brain or manager** of the chip. They dominate control logic because they cleanly map complex, sequential step-by-step algorithms into deterministic, verifiable hardware structures.

Real-World Hardware Examples

- **Serial Receiver (e.g., UART):** Moves through states like `WAIT_START`, `SAMPLE_BITS`, `CHECK_PARITY`, and `STOP_BIT`.
- **Traffic Light Controller:** Cycles through `GREEN` $\rightarrow$ `YELLOW` $\rightarrow$ `RED` based on internal timers or sensor inputs.
- **SPI/USB Controller:** Manages packet handshakes, address phases, data payload cycles, and error checking.
- **Bus Arbiter:** Evaluates request lines from multiple masters and transitions to grant states based on priority logic.

2.2 Moore vs. Mealy FSMs

The fundamental difference between these two architectures lies entirely in **how the output logic is calculated**.

Moore FSM


The outputs are a function **only of the current state**.

- If the state does not change, the output cannot change, regardless of what the inputs do.
- *Timing:* Outputs change synchronously with the clock edge because the state register updates synchronously.

Mealy FSM

The outputs are a function of **both the current state AND the current inputs**.

- If an input changes mid-cycle, the output can change immediately, even before the next clock edge arrives.
- *Timing:* Outputs can change asynchronously within a clock cycle if inputs fluctuate.



Comparison Table

Feature	Moore FSM	Mealy FSM
---------	-----------	-----------

Output Dependence	Current State only	Current State + Current Inputs
Latency	Usually 1 cycle slower (requires entering the state to assert output)	Immediate response (input paths directly bypass to output)
Glitch Risk	Safe. Outputs change cleanly on clock boundaries.	High. Input noise/glitches propagate directly to outputs.
State Count	Typically requires more states to achieve the same logic.	Typically requires fewer states .
Typical Use Cases	Dominant in RTL control paths, pipeline control.	High-speed handshaking, strict single-cycle reaction loops.

When to Choose Which

- **Choose Moore:** For the vast majority of internal RTL logic. It isolates timing paths between modules, reducing the risk of static timing violations.
- **Choose Mealy:** When a protocol strictly requires a single-cycle turnaround response to an incoming command signal.

Converting Mealy to Moore

To convert a Mealy FSM to a Moore FSM, we must eliminate the direct input dependency on the output.

1. Identify any state where an input directly forces an output.
2. Split that state into **multiple distinct states**—one for each combination of inputs that produces a unique output.
3. Shift the output generation inside those new, dedicated states.

State Encoding

State encoding defines how symbolic names (like `IDLE`) translate into literal binary bits inside the flip-flops.

1. Binary Encoding

States are assigned consecutive sequential numbers (`00`, `01`, `10`, `11`).

- **Flip-Flops Needed:** For N states. (e.g., 8 states require 3 flip-flops).

$$\log_2(N)$$

- **Pros:** Highly area-efficient.

2. One-Hot Encoding

Exactly **one** flip-flop is active (`1`) at a time, while all others are `0`. (e.g., State 0 = `0001`, State 1 = `0010`, State 2 = `0100`, State 3 = `1000`).

- **Flip-Flops Needed:** Exactly N flip-flops for N states.
- **Pros:** High speed. The next-state logic doesn't need complex binary decoding; it just checks if a specific flip-flop is high.

3. Gray Encoding

State codes are assigned such that adjacent states differ by **exactly one bit** (e.g., `00` → `01` → `11` → `10`).

- **Pros:** Minimizes dynamic power consumption and switching noise because only one bit toggles per transition.

Synthesis Tool Behavior & Trade-offs

Modern synthesis tools usually ignore explicit RTL coding values and automatically optimize encoding based on constraints:

- We can force or suggest an encoding using tool attributes directly in code: `/* synthesis syn_encoding = "one-hot" */`.

Encoding	Speed	Area (Registers)	Power	Debuggability
Binary	Slower (Deep combo logic)	Excellent (Minimal)	Medium	Harder (Must decode numbers)
One-Hot	Fastest (Shallow combo logic)	Poor (Many FFs)	Higher (More FFs toggling)	Easiest (1 bit per state)

Gray	Medium	Excellent	Lowest (Minimal toggling)	Medium
------	--------	-----------	---------------------------	--------

Transition Tables & State Diagrams

FSM design begins with a visual or tabular specification before a single line of RTL is written.

State Diagrams vs. Transition Tables

- **State Diagram:** Visual representation. **Circles** represent states. **Directed Arrows** represent transitions labeled with `Condition / Output`.
- **Transition Table:** Matrix format mapping `Current State + Inputs` $\rightarrow$ `Next State + Outputs`. It acts as a perfect functional truth table for your combinational logic block.

Verification Checklist

- **Unreachable States:** Ensure every state has at least one valid path into it from the reset state.
- **Self-Loops:** Explicitly define what happens if no transition condition is met (the FSM should explicitly hold its state: `next_state = current_state`).
- **Reset States:** Always define a deterministic starting state upon assertion of the system reset.

Two-Block FSM Style in SystemVerilog

The industry standard pattern separates sequential elements from combinational logic to ensure robust synthesis.

Code snippet

```
// Proper state type declaration using enumerated types
typedef enum logic [1:0] {
    STATE_IDLE   = 2'b00,
    STATE_READ   = 2'b01,
    STATE_WRITE  = 2'b10
} state_e;

state_e current_state, next_state;
```

Block 1: The State Register (Sequential)

Dedicated strictly to updating the state elements on the clock edge. **No logic computations happen here.**

Code snippet

```
// Example with Asynchronous Reset
always_ff @(posedge clk or negedge rst_n) begin
    if (!rst_n)
        current_state <= STATE_IDLE;
    else
        current_state <= next_state; // Non-blocking assignment
end
```

Block 2: Next-State & Output Logic (Combinational)

Calculates where to go next and what the outputs should be. Uses a `case` statement matching `current_state`.

Code snippet:

```
always_comb begin
    // Default assignments prevent latches
    next_state = current_state;
    out_signal = 1'b0;

    case (current_state)
```

```

    STATE_IDLE: begin
        if (start_input) next_state = STATE_READ;
    end
    STATE_READ: begin
        out_signal = 1'b1; // Moore output (depends purely on state)
        if (read_done) next_state = STATE_WRITE;
    end
    STATE_WRITE: begin
        if (write_done) next_state = STATE_IDLE;
    end
    default: next_state = STATE_IDLE; // Guard against illegal states
endcase
end

```

Why Two Blocks are Cleaner than One Monolithic Block

1. **Prevents Mixed Assignment Bugs:** Separates blocking (`=`) logic from non-blocking (`<=`) logic cleanly.
2. **Improves Timing Debug:** It is immediately clear which signals are purely combinational and which are registered.
3. **Tool Optimization:** Synthesis tools read this specific code pattern easily, allowing for optimal state re-encoding optimizations.

2.6 Common FSM Bugs to Avoid

⚠ Incomplete Case Statements $\rightarrow$ Inferred Latches

If a state or an output is not defined for *every possible branch* within an `always_comb` block, the synthesizer assumes the circuit must "remember" its past value. This infers a dangerous, unclocked hardware latch. Always use default assignments at the top of an `always_comb` block.

⚠ Missing Default Case

If radiation, electrical noise, or a setup violation forces the physical flip-flops into an undefined state (e.g., `2'b11` in a machine with only 3 states), the machine will lock up or misbehave unless a `default:` statement safely routes the logic back to the reset state.

⚠ Output Registered One Cycle Late

If you combine Moore outputs inside an `always_ff` block instead of an `always_comb` block, the output will change **one cycle after** the state transition occurs rather than concurrently with it. This is a common point of confusion during waveform debugging.

⚠ State Register Not Reset Properly

Forgetting to include the state register variable in the `if (!rst_n)` initialization clause causes the chip to wake up in an unknown configuration (`X`), making simulation and physical start-up non-deterministic.

⚠ Overlapping Detection Errors

In sequence detectors (e.g., searching for the pattern `1011`), failure to properly map state transitions can cause bugs when patterns overlap (e.g., `101011`). Ensure the transition out of a failed sequence lands in the correct intermediate state, rather than blindly dropping all the way back to `IDLE`.

Hardware Verification Fundamentals.

Verification vs. Validation

Digital design is completely meaningless without a rigorous process to prove that the circuits work exactly as intended.

Core Definitions

- **Verification ("Building the design right"):** Confirms that the implemented hardware description (RTL) perfectly matches the architectural specification sheet. It answers: *Did we build what we said we would build?*

- **Validation ("Building the right design"):** Confirms that the specification sheet and the manufactured chip actually satisfy the end-user's real-world requirements. It answers: *Is this chip actually useful to the customer, or did we build a perfect implementation of a flawed idea?*

Why Both are Needed & Their Timing

- **Verification** occurs **pre-silicon** (before manufacturing) using software simulators, formal tools, and emulators. Its goal is to catch mathematical and logical implementation errors early.
- **Validation** occurs **post-silicon** (after the physical chip arrives from the foundry) in a physical lab. Engineers plug the chip into actual motherboards, subject it to voltage/temperature sweeps, and boot real operating systems to check for system-level anomalies that simulation environments could never replicate.

The Cost of a Bug: The 10x Rule

Finding a bug late in the design cycle is cataclysmically expensive. The **10x Rule** states that the cost of fixing an issue multiplies by roughly ten times at each progressive milestone:

```
[RTL Simulation ($1)] → [Gate-Level Netlist ($10)] → [Tape-Out/Masks ($100k+)] → [Post-Silicon Lab ($1M+)] → [Field/Customer Recall ($100M+)]
```

If a bug is caught in RTL, a designer simply edits a line of text. If a bug escapes to post-silicon, it can require a full mask re-spin (millions of dollars) and delay time-to-market by months.

Famous Escaped Chip Bugs

- **Intel Pentium FDIV Bug (1994):** A missing entry in a lookup table inside the floating-point execution unit caused rare but critical division arithmetic errors. It cost Intel roughly \$475 million to replace defective processors.
- **Ariane 5 Rocket Crash (1996):** A software/hardware integration failure where a 64-bit floating-point variable was converted into a 16-bit signed integer, causing an arithmetic overflow that shut down the inertial reference system. The rocket veered off course and self-destructed 37 seconds after launch.

The Verification Problem

The State Space Explosion

You cannot verify a complex digital design by simply testing every possible input combination—the math makes it physically impossible.

Example: Consider a simple 32-bit adder. It has two 32-bit inputs.

Total unique input combinations = $2^{32} \times 2^{32} = 2^{64} \approx 1.84 \times 10^{19}$ combinations.

If a simulator evaluates **1 million vectors per second**, it would take over **584,000 years** to completely test that single, basic block. Add sequential state elements (flip-flops), and the possible states explode exponentially.

Coverage-Driven Verification (CDV)

Because exhaustive testing is impossible, the semiconductor industry relies on **Coverage-Driven Verification**.

- Instead of running random tests until you run out of schedule or ideas, you define objective metrics (Coverage Targets) at the start of the project.
- You stop writing and running test cases only when the simulation tools prove that 100% of those targeted structural lines and functional scenarios have been successfully exercised.

Finding Bugs vs. Proving Correctness

- **Simulation** can only show the *presence* of bugs; it can never prove their absolute *absence*. Just because a testbench passes 1 million cycles bug-free does not mean cycle 1,000,001 won't fail.
- **Formal Verification** uses mathematical static analysis (Property Checking) to prove correctness. It evaluates the design against logical properties across *all* valid state combinations simultaneously without injecting test vectors.

Testbench Anatomy

A verification testbench is an isolated software environment that completely wraps around the hardware module under test to act as its virtual world.

Core Components

1. **DUT / UUT (Device/Unit Under Test):** The raw, synthesizable RTL module instance being verified.
2. **Stimulus Generator:** Creates the high-level testing scenarios and data payloads (e.g., generating valid network packets or memory transaction streams).
3. **Driver (BFM - Bus Functional Model):** Translates abstract generator data into pin-level, cycle-accurate legal signals (toggling ready/valid handshakes, driving data buses) to feed into the DUT inputs.
4. **Monitor:** A passive observer. It attaches to the DUT input/output pins, samples the signals every clock cycle, packs them back into high-level data objects, and sends them to the checker. **Crucial rule:** A monitor must *never* drive or alter DUT signals.
5. **Checker / Scoreboard:** Contains a "Golden Reference Model" (often a high-level algorithmic description of how the chip *should* behave). It receives the actual DUT outputs captured by the monitor, compares them against the reference model's expected outputs, and flags an immediate error if a mismatch occurs.
6. **Coverage Collector:** Monitors the activity to record which bits have toggled, which state machine transitions occurred, and which protocol features were successfully exercised.

Types of Tests

Modern verification utilizes a tiered testing strategy to balance targeted precision with broad exploratory reach.

1. Directed Tests

Hand-crafted, specific inputs designed to hit an exact scenario (e.g., "force a FIFO full condition, then immediately attempt to write one more byte").

- **Pros:** Highly predictable, excellent for probing known extreme corner cases.
- **Cons:** Scalability is poor; you only find bugs you explicitly thought to look for.

2. Constrained-Random Tests (CRT)

The backbone of modern SystemVerilog verification. The generator creates thousands of random transactions, restricted only by basic legal protocol rules (e.g., "generate random packet lengths, but the length must always be between 64 and 1518 bytes").

- **Pros:** Discovers hidden, complex bugs that human designers would never think to intentionally test.
- **Cons:** Requires a complex infrastructure (Scoreboards and Monitors) to verify unexpected outputs automatically.

3. Smoke Test

A minimal, straightforward directed test run at the very beginning of a project or after an RTL modification. It verifies that basic operations work (e.g., "Can we write a register and read it back?"). If the smoke test fails, the design is obviously broken, and further testing is paused.

4. Regression

Automated suites that run dozens, hundreds, or thousands of accumulated random and directed tests across a server farm. Regressions are run periodically (often nightly or weekly) to ensure that recent RTL bug fixes or features haven't inadvertently broken older, previously working functionality (**regressions**).

The Verification Plan (VPlan)

The **Verification Plan** is a comprehensive architectural document that lists exactly *what* features of the specification need to be verified, *how* they will be checked, and *which* coverage metrics will prove completion.

Key Structural Terms

- **Feature:** A distinct, testable item described in the specification (e.g., "The design must support variable baud rates from 9600 to 115200").

- **Test Case:** A concrete implementation scenario designed to hit one or more features.
- **Priority Metrics:**
 - **P0:** Critical core paths. Must pass immediately; blocking priority.
 - **P1:** Standard operational modes and major corner cases.
 - **P2:** Rare features, performance tuning, or non-critical error handling.
- **Status Tracking:** Explicitly mapped dynamically to the regression tracking spreadsheet: `Not Started` $\rightarrow$ `In Progress` $\rightarrow$ `Pass` / `Fail`.

Why Write the VPlan First?

Writing the verification plan **before** coding the testbench prevents "confirmation bias." If you build a testbench simultaneously with code without an independent plan, you are highly likely to duplicate the same logical assumptions and blind spots that the designer put into the RTL.

Coverage Basics

Coverage provides an objective answer to the question: *"Are we done yet?"* It is split into two completely separate paradigms.

1. Code Coverage

An automated metric calculated natively by the simulator. It measures structural execution—checking if your test cases physically forced the lines of written RTL text to evaluate.

- **Statement Coverage:** Did the simulator execute this specific line of code?
- **Branch Coverage:** Did the simulation exercise both the `true` and `false` paths of an `if-else` statement or every branch of a `case` block?
- **Toggle Coverage:** Did every single bit of a bus transition cleanly from `0 -> 1` and from `1 -> 0`?
- **FSM Coverage:** Did the simulation enter every defined state, and did it exercise every valid transition link between those states?

2. Functional Coverage

A user-defined metric written explicitly in SystemVerilog code. It tracks whether the testbench successfully created the systemic scenarios required by the specification.

- **Coverpoints:** A specific variable or signal vector whose values you want to log throughout execution.
- **Bins:** User-defined storage buckets that bucket values into groups you care about.
- Code snippet:

```
covergroup PacketGroup;
  coverpoint packet_size {
    bins small  = { [64:127] };
    bins medium = { [128:1023] };
    bins jumbo  = { [1024:1518] };
  }
endgroup
```

- **Cross Coverage:** Tracks the simultaneous interaction of multiple independent coverpoints (e.g., tracking if a `Read Command` occurred exactly while the `Memory Bus is Full`).

Coverage Closure

The systematic process of analyzing your coverage reports, identifying gaps (e.g., "We have 100% code coverage, but 0% Jumbo packet functional coverage"), writing targeted tests to fill those gaps, and repeating until your targets hit 100%.

Assertions

Definition

An **Assertion** is a specialized piece of code embedded directly within the design or testbench that explicitly declares a condition that **must always be true** during active operation. If that condition ever evaluates to false, the simulator instantly halts or triggers a highly visible error message.

Immediate vs. Concurrent Assertions

- **Immediate Assertions:** Evaluated immediately at that exact instant in simulation time, behaving exactly like a basic software `if` condition statement inside procedural logic blocks (`always`).
- **Concurrent Assertions:** Evaluated over span-of-time sequences across multiple clock cycles. They are clock-driven and track complex timing-dependent properties.

Real-World Conceptual Examples

- *Protocol Check:* "If `Request` is asserted, `Grant` must be asserted within 3 clock cycles."
- *FIFO Boundary Check:* "The `full` flag must never be asserted if the internal buffer counter is strictly less than the maximum FIFO depth."
- *Reset Property:* "While `rst_n` is driven low, all output data signals must immediately clear to zero."

Why Assertions Beat Manual Testbench Checking

Manual checking in a scoreboard usually only evaluates outputs at the final boundary of a complete transaction. Assertions sit **inside** the design internals, checking properties **every single cycle**. They provide immediate visibility into a bug at the exact point of origin, preventing a internal logic failure from traveling through hundreds of clock cycles of pipelines before finally causing a noticeable error at the output pins.

SystemVerilog Assertions (SVA) — Concept Level

SystemVerilog Assertions (SVA) provide a formal, standardized way to describe and monitor the **temporal behavior** of your digital design. While standard RTL describes *how* to build the hardware, SVA describes *what the hardware must (or must not) do* over time.

Why SVA?

Traditional testbenches are "end-to-end"—they drive inputs and only check the final outputs. If a bug occurs deep inside a pipeline, it might take dozens of clock cycles to propagate to an output pin, making debugging incredibly difficult.

SVA acts as an **internal logic probe**. It sits directly inside or alongside your design blocks, continuously monitoring internal nodes and flagging violations the exact cycle they occur.

Core Concepts: Sequences and Properties

SVA is built hierarchically using two primary building blocks: **Sequences** and **Properties**.

1. Sequences (The "What" and "When")

A sequence is a boolean expression or a chain of expressions mapped across clock cycles. It defines a specific pattern of events over time.

- *Example:* "Signal `A` is high, and exactly 2 clock cycles later, signal `B` goes high."

2. Properties (The "Rules")

A property wraps sequences into complex, conditional behaviors (such as implication). It defines a formal rule or contract that the design must adhere to.

- *Example:* "If the reset is asserted, then the output must remain zero."

3. Assert (The "Action")

Simply declaring a property does not execute it. You must explicitly **assert** the property to instruct the simulator to run it, or **cover** it to track if the scenario ever occurred.

The Power of Implication (`| ->` and `| =>`)

In assertions, we often want to check a rule *only if* a specific trigger condition happens. This is called **implication**. Without implication, the simulator would constantly trigger false alarms.

SVA uses two primary implication operators:

- **Overlapping Implication (`| ->`)**: If the antecedent (trigger) is true, the consequent (consequence) is evaluated on the **same clock cycle**.
- **Non-overlapping Implication (`| =>`)**: If the antecedent is true, the consequent is evaluated on the **next clock cycle**.

Code snippet

```
// If 'req' goes high, 'ack' must go high in the EXACT SAME cycle.
property p_same_cycle;
    @(posedge clk) req | -> ack;
endproperty

// If 'req' goes high, 'ack' must go high in the NEXT cycle.
property p_next_cycle;
    @(posedge clk) req | => ack;
endproperty
```

Concurrent Assertions vs. Immediate Assertions

Feature	Immediate Assertions	Concurrent Assertions (SVA)
Trigger	Evaluated immediately like an <code>if</code> statement.	Evaluated only at specific clock boundaries.
Time Domain	Cycle-free (happens at a single instant).	Temporal (spans across multiple clock cycles).
Placement	Inside procedural blocks (<code>always</code> , <code>initial</code>).	Inside or outside procedural blocks, or in separate bind files.
Typical Use	Simple gate-level or variable bounds checking.	Protocol tracking, handshake validation, interface checking.

Benefits of Using SVA

- **Closer to the Bug**: Locates the root cause of an error immediately, eliminating hours of backtracking through wave diagrams.
- **Active in Emulation & Formal**: Unlike manual testbench checkers, SVA can be synthesized by formal tools to mathematically prove that a bug is physically impossible, or compiled into hardware emulation platforms for ultra-fast runtimes.
- **Interoperable**: SVA is standard across all industry-grade simulators (VCS, Questa, Xcelium).

Module 4: Sequence Detectors & RTL/TB Implementation

This module covers the practical implementation of Sequence Detectors using the standard two-block FSM style in SystemVerilog, alongside the fundamentals of writing a basic smoke testbench to verify your design.

Quick Review — Sequence Detector Theory

A sequence detector is a sequential circuit that outputs a pulse (usually a single-cycle high signal) when a specific pattern of incoming serial bits is matched over consecutive clock cycles.

Overlapping vs. Non-Overlapping Detection

The distinction between these two detection styles depends entirely on how the FSM treats the end of a successfully matched sequence.

Suppose we are searching for the sequence `101` :

Input Stream (w)	1	0	1	0	1	Note
------------------	---	---	---	---	---	------

Overlapping Output	0	0	1	0	1	The final 1 of the first match doubles as the first 1 of the next match.
Non-Overlapping Output	0	0	1	0	0	The FSM resets completely after a match; the next search starts fresh from the next bit.

What Changes in the State Diagram?

- **Non-Overlapping FSM:** Once the final bit of the pattern is detected, the transition from the final state always routes back to either the **IDLE** state or the single-bit start state (e.g., if the next bit doesn't match). It cannot reuse bits from the match that just occurred.
- **Overlapping FSM:** The transition out of the final state routes back to an intermediate state representing the prefix of the next potential match. For pattern **101**, after matching the second **1** (state **S_101**), if the next incoming bit is **0**, the FSM transitions to state **S_10** instead of starting over.

Moore Output Timing and Registration

In a classic Moore FSM, the output is decoded directly from the current state flip-flops using combinational logic.

However, in professional high-speed VLSI designs, pure combinational outputs are highly discouraged because they propagate gate delays directly to the next block, limiting the maximum clock frequency (F_{max}). To prevent this, we often **register the output**.

The Timing Trade-off

- **Standard Moore FSM:** The output is calculated combinational-style from the state registers. It goes high in the *exact same clock cycle* that the FSM enters the "Match" state.
- **Registered Moore FSM:** The decoded output is passed through an extra flip-flop. While this completely isolates the output timing path (breaking combinational propagation paths), it introduces a **one-clock-cycle delay**. The output pulse will appear one cycle *after* the FSM enters the match state.

Two-Block SV Implementation

This is the cleanest, most industry-accepted layout for coding synchronous state machines in SystemVerilog.

1. Declaring the State Type

Always use a strongly typed `typedef enum` with an explicit base type and width. This prevents compile-time state width mismatches and makes debugging waveforms much easier, as simulator waves will display the state names (e.g., **ST_RESET**) instead of raw numbers.

Code snippet:

```
typedef enum logic [1:0] {
    ST_IDLE   = 2'b00,
    ST_G0T_1  = 2'b01,
    ST_G0T_10 = 2'b10,
    ST_MATCH  = 2'b11
} state_e;

state_e current_state, next_state;
```

2. State Register Block (Sequential)

This block is incredibly simple. It does nothing but update the state registers on the clock edge or handle the system reset. **Do not put next-state logic calculations here.**

Code snippet:

```

always_ff @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        current_state <= ST_IDLE;
    end else begin
        current_state <= next_state; // Sequential update uses '<='
    end
end
end

```

3. Next-State & Output Block (Combinational)

This block calculates where to transition next and determines the outputs. Because it is combinational, you must assign a value to `next_state` and your outputs under **all possible conditions** to avoid inferring hardware latches.

Code snippet:

```

always_comb begin
    // 1. Default Assignments (Crucial to prevent latches!)
    next_state = current_state;
    seq_detected = 1'b0;

    // 2. State Transition and Output Logic
    case (current_state)
        ST_IDLE: begin
            if (din) next_state = ST_GOT_1;
        end

        ST_GOT_1: begin
            if (!din) next_state = ST_GOT_10;
            // if din is 1, we stay in ST_GOT_1 (re-using the 1)
        end

        ST_GOT_10: begin
            if (din) next_state = ST_MATCH;
            else      next_state = ST_IDLE;
        end

        ST_MATCH: begin
            seq_detected = 1'b1; // Moore output: depends purely on being in ST_MATCH
            // Overlapping logic for sequence '101':
            // The current active bit is the '1' that completed the previous match.
            if (!din) next_state = ST_GOT_10;
            else      next_state = ST_GOT_1;
        end

        default: next_state = ST_IDLE;
    endcase
end

```

Reset Handling: Synchronous vs. Asynchronous

Reset Type	Code Syntax	Pros	Cons
Asynchronous	<code>always_ff @(posedge clk or negedge rst_n)</code>	Works even if the clock is completely dead; does not require a clock edge to put the system in a safe state. Saves logic gates.	Susceptible to glitching on the reset line. Requires special synchronization circuits during deassertion to prevent metastability.
Synchronous	<code>always_ff @(posedge clk)</code>	Clean, fully synchronous design. Synthesizer can easily optimize timing. Inherently filtered against narrow glitches.	Requires an active clock signal to reset the system. If the clock is missing, the circuit cannot be initialized.

Design Practice: Modern ASIC and FPGA designs heavily favor **Asynchronous Assert, Synchronous Deassert** resets. This uses an asynchronous reset structure for immediate system clearing but applies a "reset synchronizer" circuit to release the reset cleanly on a clock edge.

Smoke Testbench Basics

A smoke testbench is a non-synthesizable file designed to wrap around your module, drive some basic signals, and let you check if the logic is fundamentally sound.

Minimal Smoke Testbench Template

Code snippet:

```
`timescale 1ns/1ps

module tb_seq_detector;

    // 1. Declare local signals to connect to the DUT
    logic clk;
    logic rst_n;
    logic din;
    logic seq_detected;

    // 2. Instantiate the Device Under Test (DUT)
    seq_detector u_dut (
        .clk      (clk),
        .rst_n    (rst_n),
        .din      (din),
        .seq_detected (seq_detected)
    );

    // 3. Clock Generator Block (Runs infinitely)
    initial begin
        clk = 0;
        forever #5 clk = ~clk; // 10ns clock period (100 MHz)
    end

    // 4. Stimulus Block
    initial begin
        // Initialize inputs
        din  = 0;
        rst_n = 0;

        // Hold reset for 2 clock cycles
        #20;
        rst_n = 1; // Release reset

        // Drive input sequence '101'
        @(posedge clk); din = 1;
        @(posedge clk); din = 0;
        @(posedge clk); din = 1; // At the next clock, match should occur

        // Wait a few cycles and finish
        repeat (5) @(posedge clk);
        $finish; // Safely stop the simulator
    end

    // 5. Monitoring Signals in the Terminal
    initial begin
```

```

        $monitor("Time=%0t | rst_n=%b | din=%b | State=%s | Detected=%b",
                $time, rst_n, din, u_dut.current_state.name(), seq_detected);
    end

endmodule

```

System Tasks: `$display` vs. `$monitor`

- `$display`: Actively prints the specified string **once** at the exact moment the line is executed in the procedural code block (highly analogous to `printf` in C).
- `$monitor`: A background process that prints its string **automatically whenever any of its arguments change value**. You should only instantiate a single `$monitor` block inside a testbench to prevent terminal spam.

Verifying Outputs "By Eye"

When debugging your sequence detector design using simulator waves (such as GTKWave, ModelSim, or Verdi) or looking at terminal output logs, use this quick checklist:

1. **Locate the Reset Window:** Confirm that while `rst_n` is `0`, the internal state registers are firmly resolved to `ST_IDLE` (no unknown `x` states).
2. **Verify Clock Alignment:** Ensure your input data `din` transitions *away* from the active clock edge (typically on the falling edge or a small delay after the rising edge) to avoid simulation setup-and-hold anomalies.
3. **Trace State History:** Ensure that as the input bits line up with your target sequence on the clock edges, the internal state shifts sequentially through `ST_GOT_1` → `ST_GOT_10` → `ST_MATCH`.
4. **Confirm Output Pulse Width:** Verify that `seq_detected` stays high for exactly one full clock cycle when a match is found and returns to `0` on the subsequent cycle (unless another overlapping match occurs immediately).

Module 5: FSM + Datapath (FSDP) & Output Encoding

This module covers the FSM+Datapath (FSDP) design pattern—a cornerstone of professional digital design used to handle complex timing, algorithms, and physical control structures like traffic light controllers.

5.1 FSM + Datapath (FSDP) Pattern

In basic FSMs, state transitions occur on every single clock cycle based purely on immediate input signals. However, real-world controller states often need to persist for specific, variable, or long durations.

Why a Pure FSM is Insufficient for Durations

If a state (e.g., `MAIN_GREEN`) needs to last for 10,000 clock cycles, implementing this in a **pure FSM** is highly inefficient:

- **State Explosion:** You would need to manually or algorithmically declare 10,000 distinct states (`GREEN_1`, `GREEN_2`, ..., `GREEN_10000`) to act as a cycle counter.
- **Logic Bloat:** The resulting synthesis would generate a massive, highly congested combinational transition network, wasting layout area and degrading timing performance.

The Solution: Decouple Control from Arithmetic

To solve this, we split the system into two distinct architectural blocks:

1. **The Control Path (FSM):** A simple state machine that manages high-level logical modes (e.g., `GREEN`, `YELLOW`, `RED`).
2. **The Datapath (Counter):** A dedicated arithmetic unit (registers and decrementers) that handles the actual numerical counting.

The Feedback Loop Handshake

The interaction between the FSM and the Datapath relies on a clean, 2-wire handshake:

- **Control Signal (FSM → Datapath):** The FSM asserts a `load` or `start` signal to initialize the counter with a target duration value when entering a state.

- **Status Signal (Datapath → FSM):** The counter counts down independently. Once it reaches zero, it asserts a `timer_done` (or `zero`) status flag.
- **Transition Rule:** The FSM sits in its active state and **ignores** normal transition inputs *until* the Datapath asserts `timer_done`.

Parameterizing State Durations

Never hardcode literal numbers (e.g., `count <= 32'd50000`) inside your RTL logic. Instead, use parameters:

```
parameter CLK_FREQ_HZ    = 50_000_000;
parameter SEC_5_CYCLES  = CLK_FREQ_HZ * 5;
parameter SEC_1_CYCLES  = CLK_FREQ_HZ * 1;
```

Why Hardcoding is Bad Practice:

1. **Unreadable Code:** Future designers will have no idea what a magic number like `50000` means.
2. **Zero Portability:** If the system clock speed changes from 50 MHz to 100 MHz, you must manually hunt down and recalculate every single number in the codebase.
3. **Simulation Nightmare:** Simulating a real-world 10-second delay at 100 MHz requires simulating 1,000,000,000 clock cycles, which can freeze or drastically slow down your simulator.

5.2 Counter Datapath in SystemVerilog

A robust down-counter serves as the standard timer engine for our FSM.

Load vs. Decrement Logic

The counter must prioritize loading a new value (when transitioning states) over its default behavior of decrementing on each clock cycle.

```
module down_counter #(
    parameter WIDTH = 32
) (
    input  logic      clk,
    input  logic      rst_n,
    input  logic      load,
    input  logic [WIDTH-1:0] load_val,
    output logic      timer_done
);

    logic [WIDTH-1:0] count;

    // Sequential counting logic
    always_ff @(posedge clk or negedge rst_n) begin
        if (!rst_n) begin
            count <= '0;
        end else if (load) begin
            count <= load_val - 1'b1; // Load and immediately account for 1 cycle delay
        end else if (count > '0) begin
            count <= count - 1'b1;    // Active countdown
        end
    end

    // Status output: asserted when counter hits zero
    assign timer_done = (count == '0);

endmodule
```

Parameter Override for Verification (Testability)

To test long-duration FSMs without waiting hours for simulation runs, you can override the duration parameters during module instantiation in your testbench:

```
// In actual hardware, standard green lasts 5,000,000 cycles
traffic_controller #(
    .GREEN_DURATION(5_000_000)
) u_dut_hardware (...);

// In your testbench, override it to 5 cycles for fast verification!
traffic_controller #(
    .GREEN_DURATION(5)
) u_dut_sim (...);
```

5.3 Output Encoding

For controllers driving safety-critical outputs (like traffic lights), how you structure your output values is just as important as how you structure your state transitions.

Driving Multi-Bit Outputs

Consider a standard two-road intersection (Main Road and Side Road). Each road has three lights (Red, Yellow, Green), requiring a total of 6 control outputs.

```
// Grouping outputs as a single status vector: [M_R, M_Y, M_G, S_R, S_Y, S_G]
logic [5:0] lights;
```

One-Hot Output Style

To prevent catastrophic electric failures (like both roads accidentally displaying Green at the same time due to a logic hazard), map the outputs directly to your state names using **One-Hot Encoding** principles.

FSM State	Main Road Lights	Side Road Lights	6-bit Vector Value
ST_MAIN_GREEN	Green	Red	6'b001_100
ST_MAIN_YELLOW	Yellow	Red	6'b010_100
ST_ALL_RED	Red	Red	6'b100_100
ST_SIDE_GREEN	Red	Green	6'b100_001
ST_SIDE_YELLOW	Red	Yellow	6'b100_010

Managing the Yellow State & The "All-Red" Phase

A common rookie mistake is transitioning directly from one road's Yellow state straight to the other road's Green state:

Main Yellow → Side Green

The Danger of Immediate Transition

In physical hardware, there are propagation delays and physical light bulb turn-off latencies. If you switch instantly, there may be a fraction of a second where drivers on both roads see green/yellow overlapping, leading to accidents.

The Safe Design Standard:

Always implement an intermediate **All-Red Phase** (`ST_ALL_RED`) between transitions:

Main Yellow → All Red (1-2 seconds) → Side Green

This ensures that the intersection is entirely cleared of vehicles before the perpendicular traffic flow is given a green light.

Single-Point FSM + Datapath Design Template

Below is a complete, clean, self-contained template combining an FSM and a parameterized counter datapath in a single SystemVerilog module:

```
module traffic_light_controller #(
    parameter CLK_FREQ = 50_000_000,
    parameter T_GREEN  = 30 * CLK_FREQ, // 30 seconds
    parameter T_YELLOW = 5  * CLK_FREQ, // 5 seconds
    parameter T_RED    = 2  * CLK_FREQ  // 2 seconds (All-Red clearance)
) (
    input  logic clk,
    input  logic rst_n,
    output logic [2:0] main_lights, // [Red, Yellow, Green]
    output logic [2:0] side_lights  // [Red, Yellow, Green]
);

// 1. State Declarations
typedef enum logic [2:0] {
    ST_M_GREEN  = 3'b000,
    ST_M_YELLOW = 3'b001,
    ST_ALL_RED1 = 3'b010,
    ST_S_GREEN  = 3'b011,
    ST_S_YELLOW = 3'b100,
    ST_ALL_RED2 = 3'b101
} state_e;

state_e current_state, next_state;

// 2. Timer Datapath Signals
logic [31:0] timer_count;
logic [31:0] timer_load_val;
logic        timer_load;
logic        timer_done;

// 3. Sequential State Register
always_ff @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        current_state <= ST_ALL_RED1;
    end else begin
        current_state <= next_state;
    end
end

// 4. Integrated Down-Counter Datapath
always_ff @(posedge clk or negedge rst_n) begin
    if (!rst_n) begin
        timer_count <= '0;
    end else if (timer_load) begin
        timer_count <= timer_load_val - 1'b1;
    end else if (timer_count > '0) begin
        timer_count <= timer_count - 1'b1;
    end
end

assign timer_done = (timer_count == '0);

// 5. Combinational Next-State & Control Logic
always_comb begin
```

```

// Default assignments to prevent latches
next_state      = current_state;
timer_load      = 1'b0;
timer_load_val = T_RED;

case (current_state)
    ST_M_GREEN: begin
        timer_load_val = T_GREEN;
        if (timer_done) begin
            next_state = ST_M_YELLOW;
            timer_load = 1'b1; // Load yellow duration for next state
        end
    end

    ST_M_YELLOW: begin
        timer_load_val = T_YELLOW;
        if (timer_done) begin
            next_state = ST_ALL_RED1;
            timer_load = 1'b1;
        end
    end

    ST_ALL_RED1: begin
        timer_load_val = T_RED;
        if (timer_done) begin
            next_state = ST_S_GREEN;
            timer_load = 1'b1;
        end
    end

    ST_S_GREEN: begin
        timer_load_val = T_GREEN;
        if (timer_done) begin
            next_state = ST_S_YELLOW;
            timer_load = 1'b1;
        end
    end

    ST_S_YELLOW: begin
        timer_load_val = T_YELLOW;
        if (timer_done) begin
            next_state = ST_ALL_RED2;
            timer_load = 1'b1;
        end
    end

    ST_ALL_RED2: begin
        timer_load_val = T_RED;
        if (timer_done) begin
            next_state = ST_M_GREEN;
            timer_load = 1'b1;
        end
    end

    default: next_state = ST_ALL_RED1;
endcase
end

```



```
// 6. Outputs mapped directly from state (Clean, hazard-free decoding)
always_comb begin
    // Default outputs: Red/Red for safety
    main_lights = 3'b100;
    side_lights = 3'b100;

    case (current_state)
        ST_M_GREEN: begin main_lights = 3'b001; side_lights = 3'b100; end
        ST_M_YELLOW: begin main_lights = 3'b010; side_lights = 3'b100; end
        ST_ALL_RED1: begin main_lights = 3'b100; side_lights = 3'b100; end
        ST_S_GREEN: begin main_lights = 3'b100; side_lights = 3'b001; end
        ST_S_YELLOW: begin main_lights = 3'b100; side_lights = 3'b010; end
        ST_ALL_RED2: begin main_lights = 3'b100; side_lights = 3'b100; end
        default: begin main_lights = 3'b100; side_lights = 3'b100; end
    endcase
end

endmodule
```

Module 6: Synchronous FIFO Design & Pointer Architecture

This module covers the core fundamentals, architecture, flag generation, and pointer mathematics behind a **Synchronous FIFO (First-In, First-Out)** buffer.

6.1 What is a FIFO and Why It Exists

A **FIFO** is a dual-port storage element used to pass data between different parts of a digital system. As the name implies, it operates on a strict queuing principle: the first piece of data written into the buffer is the very first piece of data read out of it.

Core Applications in SoC Design

FIFOs are ubiquitous in System-on-Chip (SoC) architectures, primarily serving three functions:

- **Rate Matching (Elastic Buffering):** Bridges the gap between a high-speed producer and a low-speed consumer (or vice versa). For example, a high-speed PCIe core can burst data into a FIFO, while a slower internal peripheral reads it out at a manageable pace.
- **Bursty Traffic Management:** Temporarily stores incoming network packets or memory blocks when the downstream receiver is busy, preventing data loss during temporary bottlenecks.
- **Clock Domain Crossing (CDC):** Acts as a safe boundary for passing data between entirely asynchronous clock domains (Note: This requires an **Asynchronous FIFO**, which uses specialized synchronization logic like Gray code pointers, unlike the simpler Synchronous FIFO covered here).

Synchronous vs. Asynchronous FIFOs

- **Synchronous FIFO:** Both the write interface and the read interface operate on the **exact same clock domain**. The pointer logic and flag generation are straightforward because there are no setup/hold hazards between the read and write sides.
- **Asynchronous FIFO:** The write and read interfaces run on completely independent, unrelated clocks, requiring multi-stage flip-flop synchronizers to prevent metastability.

6.2 FIFO Architecture

A basic Synchronous FIFO consists of a storage array managed by two independent tracking pointers.

1. The Memory Array

The physical storage is typically implemented as a multi-word register file or an embedded Block RAM (SRAM). In SystemVerilog, this is represented as a 2D array:

Code snippet:

```
parameter WIDTH = 8;
parameter DEPTH = 16;
logic [WIDTH-1:0] mem [0:DEPTH-1]; // DEPTH words, each WIDTH bits wide
```

2. The Write Pointer (`wr_ptr`)

Points strictly to the specific memory address where the *next incoming data word* will be written. It increments by 1 after every successful write operation.

3. The Read Pointer (`rd_ptr`)

Points strictly to the memory address of the *oldest valid data word* currently waiting to be read. It increments by 1 after every successful read operation.

4. Pointer Wraparound

Because memory arrays are finite, the pointers operate in a circular fashion. When a pointer reaches the last address (`DEPTH-1`), the next increment resets it back to address `0`.

6.3 Full & Empty Flag Logic

Flag generation is the most safety-critical portion of FIFO design. It prevents the system from corrupting data or reading garbage values.

Flag Rules

- **empty Status:** Asserted high when `count == 0`. It signals that there is no valid data left to read.
- **full Status:** Asserted high when `count == DEPTH`. It signals that the memory array is entirely filled.

The Operation Matrix

The counter register tracking the FIFO's status must handle multiple concurrent events at every rising clock edge:

Write Request (wr)	Read Request (rd)	Condition / Action	Effect on count
1	0	Successful Write (only if <code>!full</code>)	Increments (<code>count <= count + 1</code>)
0	1	Successful Read (only if <code>!empty</code>)	Decrements (<code>count <= count - 1</code>)
1	1	Simultaneous Read and Write	Unchanged (<code>count <= count</code>)
0	0	Idle	Unchanged

Boundary Overruns (Hardware Protection)

Professional RTL must contain hardware guards to ignore illegal requests:

- **Write When Full:** The FIFO must actively block internal write operations if `full == 1`. If not blocked, new data will overwrite existing unread data (**Overflow**), causing catastrophic corruption.
- **Read When Empty:** The FIFO must actively block internal read operations if `empty == 1`. If not blocked, the FIFO outputs old, invalid data (**Underflow**).

6.4 Pointer Arithmetic & Flag Ambiguity

A common design hurdle when managing circular buffers is detecting the exact boundary between a completely empty FIFO and a completely full FIFO.

The Ambiguity Problem

If we try to determine full and empty flags *solely* by comparing the raw pointers (`wr_ptr == rd_ptr`), you run into a mathematical paradox:

- When the FIFO is completely **empty**, `wr_ptr` equals `rd_ptr`.
- When the FIFO is completely **full** and has wrapped around the memory array, `wr_ptr` *also* equals `rd_ptr`.

Solving the Ambiguity: The Two Approaches

Approach A: The Count Register (Simpler, Used Below)

You keep a dedicated register (`count`) that tracks the total number of items currently sitting inside the FIFO.

- If `wr_ptr == rd_ptr` and `count == 0` $\rightarrow$ FIFO is empty.
- If `wr_ptr == rd_ptr` and `count == DEPTH` $\rightarrow$ FIFO is full.

Approach B: The Extra Pointer Bit (N+1 Bits)

Instead of sizing your pointer variables to address exactly `DEPTH` locations (e.g., a 3-bit pointer for an 8-depth FIFO), you expand the pointers by **one extra bit** (making them 4 bits wide).


- The lower N bits act as the actual memory index.
- The most significant bit (MSB) acts as a "wraparound flag."
- If the lower bits match but the MSBs are **equal**, the FIFO is **empty**.
- If the lower bits match but the MSBs are **different**, the FIFO is **full**.

Module 7: RTL Specification & Design Documentation

This module covers the principles of writing high-quality design specifications, mapping out functional requirements, and detailing boundary behaviors to ensure a flawless handoff between design and verification teams.

7.1 What is an RTL Specification Document?

An **RTL Specification (Spec)** is the authoritative, golden reference document for a hardware design block. It bridges the gap between high-level architectural concepts and physical implementation.



The Spec as a "Contract"

The spec is a strict contract signed between the **Design Engineer** and the **Design Verification (DV) Engineer**. It defines the boundaries of correct behavior.

- If a behavior is in the spec, the design *must* implement it, and the testbench *must* check it.
- If a behavior is not in the spec, it does not exist as far as the project scope is concerned.

Verification from the Spec Alone

A core tenet of professional ASIC development is that **a DV engineer must be able to construct the entire verification environment and write all test cases using only the spec document**, without looking at a single line of the designer's RTL code.

- This prevents **confirmation bias**. If the verifier reads the RTL code to understand how the block works, they will naturally duplicate any logical errors or missing edge cases the designer overlooked in their testbench code.

Resolving Ambiguities

When a spec is poorly written or vague, the designer and verifier will inevitably interpret the functional requirements differently.

- **The Symptom:** During simulations, a test fails because the testbench scoreboard expects Output A while the RTL produces Output B.
- **The Resolution:** These conflicts must not be settled by "whoever wrote the code first." Instead, the design and verification engineers must sit down, clarify the functional intent, update the spec document first, and then modify both the RTL and testbench to match the newly clarified spec.

Timing: When is the Spec Written?

The spec is a living document, but it must be established at a specific milestone:

- **Write it after** the high-level architecture and microarchitecture of the block are structurally stable.
- **Write it before** verification planning or testbench development begins. Writing a testbench without a stable spec is like building a house without a blueprint.

7.2 Sections of a Good RTL Spec

A professional-grade RTL specification contains the following structure:

1. Overview

A high-level introduction explaining what the block is, why it exists, where it sits inside the larger System-on-Chip (SoC) hierarchy, and key architectural constraints (such as target frequency or power limits).

2. Key Features

A concise, bulleted list detailing the explicit capabilities of the block.

Example (SPI Master):

- Supports Motorola SPI modes 0, 1, 2, and 3.
- Parameterizable clock prescaler (even divisions from 2 to 256).
- Hardware-controlled Slave Select (SS) lines up to 4 slaves.

3. Interface Table

An exhaustive table detailing every single input and output pin on the block boundary (see Section 7.3).

4. Microarchitecture

An explanation of the internal structural design. It includes block diagrams, clock/reset structures, descriptions of registers, internal data paths, and a list of design parameters.

5. FSM Specification

An explicit breakdown of any finite state machines controlling the module, including state mappings, transition rules, and diagrams (see Section 7.4).

6. Timing Diagrams

Key waveforms showing how input and output signals must behave relative to the clock. This is critical for showing handshake sequences (like ready/valid interfaces) and bus cycles.

7. Corner Cases

A targeted registry of boundary scenarios and dangerous operational states, documenting exactly how the hardware safely recovers from them (see Section 7.5).

8. RTL Code

The final, verified SystemVerilog code of the module is appended at the very end of the document (or linked via a version-control repository path) for reference.

7.3 Writing a Good Interface Table

The interface table is the most heavily referenced section of a specification. If a pin width or active level is documented incorrectly, compilation errors and physical connection mismatches will occur downstream.

Interface Table Design Rules

1. **Never omit a port.** Every single pin, including test ports, clocks, and resets, must have its own row.
2. **Be precise with widths.** Do not just write `data` —write `data[7:0]` or `data[WIDTH-1:0]`.
3. **Document the clock domain.** If your design is multi-clock, explicitly declare which clock controls each signal.
4. **Define active levels.** Specify whether control signals are active-high or active-low (e.g., write `rst_n` as active-low).
5. **Write behavioral descriptions.** Do not describe a port with its own name. Write *when* it is asserted and *what hardware action it causes*.

Example: Sync FIFO Interface Table

Port Name	Direction	Width	Clock Domain	Active Level	Description
<code>clk</code>	Input	1	-	Rising Edge	Main system clock. All synchronous elements update on this edge.
<code>rst_n</code>	Input	1	<code>clk</code>	Low	Asynchronous system reset. When asserted, clears pointers and count to 0 immediately.
<code>wr_en</code>	Input	1	<code>clk</code>	High	Write enable. When asserted, writes <code>din</code> into the storage array at the current write pointer. Ignored if <code>full</code> is active.
<code>din</code>	Input	<code>WIDTH</code>	<code>clk</code>	-	Input data bus to be stored in the FIFO memory.
<code>full</code>	Output	1	<code>clk</code>	High	FIFO full flag. Asserted when the internal entry counter equals the defined <code>DEPTH</code> . Blocks further writes.
<code>rd_en</code>	Input	1	<code>clk</code>	High	Read enable. When asserted, drives the oldest data word onto <code>dout</code> and advances the read pointer. Ignored if <code>empty</code> is active.
<code>dout</code>	Output	<code>WIDTH</code>	<code>clk</code>	-	Output data bus displaying the oldest unread FIFO entry.
<code>empty</code>	Output	1	<code>clk</code>	High	FIFO empty flag. Asserted when the internal entry counter equals 0. Blocks further reads.

7.4 Writing a Good FSM Spec

An FSM spec must remove all guesswork regarding control sequences.

1. State Encoding Table

Maps symbolic state names to their literal, synthesized hardware values.

State Name	Encoding (Binary)	Functional Description
<code>ST_IDLE</code>	<code>2'b00</code>	Wait state. Core is inactive, awaiting the <code>start</code> trigger.
<code>ST_READ</code>	<code>2'b01</code>	Read phase. Drives address lines and captures RAM data.
<code>ST_WRITE</code>	<code>2'b10</code>	Write phase. Dispatches accumulated data payload to output.

2. Transition Table


Defines the state transition matrix cleanly.

Current State	Transition Condition	Next State	Outputs Asserted
<code>ST_IDLE</code>	<code>start == 1</code>	<code>ST_READ</code>	None
<code>ST_IDLE</code>	<code>start == 0</code>	<code>ST_IDLE</code>	None
<code>ST_READ</code>	<code>read_done == 0</code>	<code>ST_READ</code>	<code>mem_read_en = 1</code>
<code>ST_READ</code>	<code>read_done == 1</code>	<code>ST_WRITE</code>	<code>mem_read_en = 1</code>
<code>ST_WRITE</code>	<code>write_done == 0</code>	<code>ST_WRITE</code>	<code>mem_write_en = 1</code>

ST_WRITE	write_done == 1	ST_IDLE	mem_write_en = 1, done_pulse = 1
----------	-----------------	---------	----------------------------------

3. ASCII State Diagram

An embedded text-based diagram allowing quick visual reference directly in the markdown, text file, or code comments.



7.5 Writing Good Corner Cases

Corner cases are boundary scenarios where logic is most likely to fail (such as resource exhaustion, timing conflicts, or illegal inputs). A good spec documents these explicitly so both designers and verifiers can target them.

Finding Corner Cases

To identify potential structural failure points, analyze:

- **The extremes of parameters:** What happens when `DEPTH` is set to its absolute minimum value (e.g., 1 or 2)?
- **Signal combinations:** What happens if control signals assert simultaneously (such as a read and write request on the exact same clock edge)?
- **Asynchronous boundaries:** What happens if a reset occurs in the middle of an active multicyle transaction?

Example: Sync FIFO Corner Cases Registry

Scenario	Expected Safe Behavior	How it is Handled in RTL
Write Attempt When Full	The write operation must be ignored. The write pointer must not advance, and existing data must remain uncorrupted.	The internal write command checks the full boundary condition: <code>write_valid = wr_en && !full;</code>
Read Attempt When Empty	The read operation must be ignored. The read pointer must not advance, and the output bus <code>dout</code> must not output invalid data.	The internal read command checks the empty boundary condition: <code>read_valid = rd_en && !empty;</code>
Simultaneous Read and Write (FIFO Empty)	The write should succeed, and the data should be available for reading. The count register must remain at 0 (write and read offset).	Handled by updating the write pointer, loading the input directly to memory, and keeping <code>count</code> constant inside the FSM.
Simultaneous Read and Write (FIFO Full)	The read should succeed (making a slot available), and the write should succeed in that newly freed slot. The count register remains at <code>DEPTH</code> .	The internal write and read operations evaluate concurrently, bypassing flag locks because both signals assert at once.
Reset Mid-Transaction	The FIFO must instantly clear all storage. Pointers and counters must drop to 0 within the same cycle, discarding incomplete frames.	The asynchronous <code>rst_n</code> pin triggers an immediate reset of pointers and status values inside the <code>always_ff</code> block.

Module 8: Design Consistency & Spec Refinement

This module focuses on maintaining absolute structural alignment across multiple design blocks, enforcing strict naming conventions, and utilizing the specification process as an RTL optimization tool.

8.1 Consistency Across Multiple Designs

When moving from a single module to managing an entire System-on-Chip (SoC) or multi-block pipeline, inconsistency becomes a significant source of integration bugs. If Block A defines a reset as active-high and Block B defines it as active-low, the top-level integration will fail.

Using a Uniform Specification Template

Every module specification within a project must follow the exact same structural template. Whether you are documenting a simple sequence detector, a complex traffic light controller, or a synchronous FIFO, the document layout should be identical:

Overview → Key Features → Interface Table → Microarchitecture → FSM/Logic Spec → Corner Cases

Why Uniform Templates Matter:

- **Faster DV Onboarding:** Design Verification engineers can skim multiple specs and immediately locate the interface tables or corner case registries without hunting through different layouts.
 - **Automated Parsing:** Modern VLSI teams often use scripts (Python/Perl) to parse specification markdown or Excel tables to automatically generate testbench wrappers, port connections, or documentation registers. Standardizing the layout makes this automation possible.
-

Strict Naming Conventions

The names used in your functional specification must match the literal SystemVerilog RTL code down to the exact case, character, and bit-width.

Recommended Industry Naming Rules:

- **Clock & Reset:** Always use standardized names like `clk` and `rst_n` (the `_n` suffix explicitly denotes an active-low signal).
- **Ports:** Use consistent lower-case or snake_case for input/output ports (e.g., `din`, `dout`, `wr_en`, `seq_detected`).
- **Parameters:** Use upper-case for compile-time constants to distinguish them from wires or variables (e.g., `WIDTH`, `DEPTH`, `CLK_FREQ`).
- **Enumerated States:** Prefix your state names with `ST_` or `STATE_` inside your enums (e.g., `ST_IDLE`, `ST_MATCH`).

✗ Bad Spec vs. RTL Disconnect:

- *Spec says:* An output flag named `Seq_Detected` indicates a match.
 - *RTL code contains:* `logic seq_det_o;`
 - *Result:* Compilation breaks, and automated verification binding tools fail.
-

The Spec-Writing Feedback Loop

Writing a specification is not just a passive paperwork exercise; it is an active design review method. Often, when you force yourself to document an interface table or map out an explicit state transition matrix, you will uncover hidden assumptions or logical flaws in your code.

Common Issues Uncovered by Spec Writing:

1. **Asymmetric Data Widths:** You might realize an internal accumulation register is only 8 bits wide, but under certain corner cases, it could overflow. Writing the spec forces you to explicitly define these widths.
2. **Ambiguous Reset Values:** While filling out the state encoding or register tables, you may notice that certain control registers lack a clear default state when `rst_n` is asserted.

3. **Mismatched Active Levels:** You might find that some sub-modules evaluate enabling signals on a high level (`enable == 1`), while adjacent blocks expect a low level, requiring you to refactor your RTL for consistency.


Golden Rule: When documentation reveals a gap, naming conflict, or architectural oversight, **always fix the specification first**, treat it as the updated contract, and then modify the RTL code to conform to it. Never let your code and your documentation drift apart.

Module 9: The Verification Plan (VPlan) & Requirement Traceability

This module covers the methodologies for constructing a professional Verification Plan (VPlan), mapping and identifying features, structuring test suites, and building a robust traceability matrix.

9.1 The Verification Plan as a Document

The **Verification Plan (VPlan)** is the master blueprint that guides the entire verification process. It translates the raw requirements of the RTL Specification into a concrete, executable testing strategy.



The Three Critical Questions

A complete verification plan must explicitly answer three foundational questions:

1. **What are we verifying?** (The complete list of features, protocols, and interfaces).
2. **How are we verifying it?** (Testbench structures, directed vs. constrained-random strategies, assertion strategies, or formal analysis).
3. **How do we know when we are done?** (Objective exit criteria, such as coverage percentages and bug closure metrics).

Stakeholders of the VPlan

- **The Author:** Typically written by the **Design Verification (DV) Lead** (or the DV engineer assigned to the block), with active consultation from the RTL designer to identify hidden design complexities.
- **The Users:**
- **DV Engineers:** To guide the design of the testbench, stimulus generators, monitors, and assertions.
- **Project Managers:** To track milestones, allocate engineering resources, and measure schedule progress.
- **Sign-off Reviewers:** During tape-out reviews, this document serves as the primary evidence proving the chip was verified thoroughly.

A Living Document

The VPlan is never locked away and forgotten. As design bugs are uncovered, RTL changes are made, or additional architectural specifications are added, the VPlan is continuously refined and updated.

9.2 Feature Identification

Before you can write a test case, you must break down the RTL specification into a granular list of testable targets called **Features**.

Defining a Feature

A feature is a named, independently testable behavior of the Device Under Test (DUT). It represents a single functional promise made by the specification.

Sources of Features

Features do not just come from the primary description. They are extracted from:

- The **Interface Table** (e.g., "The reset signal must instantly clear all outputs").
- The **FSM Specification** (e.g., "Must transition to error state if invalid sequence arrives").
- The **Corner Case Registry** (e.g., "Simultaneous read and write when FIFO is full must succeed without losing data").
- **Industry Protocols** (e.g., AMBA AXI handshaking rules).

Feature Granularity

Getting the detail level right is a key skill for a verification engineer:

- **✗ Too Coarse (Broad):** `write_operation` — This is too generic. It is impossible to link this to a single verification test or a precise coverage metric.
- **✗ Too Fine (Specific):** `wr_ptr_increments_on_cycle_3` — This is an implementation detail. If the RTL designer shifts the pipeline latency tomorrow, this feature description breaks even though the functionality is identical.
- **Just Right:** `fifo_write_stores_data_correctly` or `write_pointer_wraps_around_at_depth` — This is functional, design-agnostic, and completely testable.

9.3 Test Categories

To ensure a highly resilient chip, your VPlan must organize test cases into specific functional categories to target distinct operational areas.

1. Reset / Initialization Tests

Verify that the chip exits the reset state consistently.

- *Example:* Ensuring that all internal registers, state machines, and output lines clear to safe, known default values within one clock cycle of `rst_n` going low.

2. Functional Tests ("The Happy Path")

Verify standard, expected operational scenarios.

- *Example:* For a FIFO, writing 5 items and reading 5 items back to ensure they match.

3. Corner Case / Boundary Tests

Expose bugs hidden at the limits of the physical parameters.

- *Example:* For a FIFO, checking what happens when it is exactly 1 slot away from being full, or testing a sequence detector with a continuous stream of overlapping targets.

4. Negative Tests (Error Handling)

Verify that the DUT is robust enough to reject illegal inputs or abnormal protocol behaviors without hanging or corrupting internally.

- *Example:* Attempting to read from an empty FIFO, or driving data inputs while `rst_n` is asserted.

5. Stress / Consecutive Tests

Verify that sustained, highly congested, or high-speed operations do not cause internal memory leaks, buffer overflows, or progressive timing drifts.

- *Example:* Running random writes and reads continuously for millions of cycles at maximum throughput.

9.4 Requirement IDs (Req-IDs)

Every test scenario in a VPlan is assigned a permanent, unique identifier called a **Req-ID**.

Structuring a Req-ID

A Req-ID utilizes a structured prefix to represent its operational category, followed by a sequential index number:

[Category Prefix]_[Sequential ID]

- **RST_001** : Reset initialization test.
- **FUNC_003** : Core functional test.
- **CORNER_002** : Corner case boundaries test.
- **NEG_001** : Illegal input error handling test.
- **STRESS_004** : Long-term execution test.

Why Req-IDs Matter: Traceability

Req-IDs act as the primary indexing key across your entire project infrastructure. They allow you to build an explicit trail from the initial specification down to the final verification run:

RTL Spec Section 5.2 → VPlan Req-ID FUNC_003 → Test Name: `test_fifo_read_write.sv` → Coverage

The Rule of Permanent IDs

Once a Req-ID is defined in the VPlan, **it must never change or be reused**.

- If a specific test scenario becomes obsolete due to a hardware feature removal, that ID is officially **retired** (e.g., marked as **RETIRED** or **DEPRECATED** in your tracking document).
- Do not slide downstream IDs up to fill the gap. This prevents historical regression metrics or bug ticket references from pointing to completely different test scenarios.

9.5 Entry and Exit Criteria

To keep verification projects moving forward predictably, you must define strict boundaries to determine when verification is legally allowed to start and when it is complete.

Entry Criteria

The baseline conditions that must be satisfied before testbench construction and active simulation runs can begin:

- The **RTL Specification** is finalized and approved by stakeholders.
- The **Verification Plan** is complete and reviewed.
- The **Testbench Architecture** is mapped out, and necessary simulator licenses are secured.
- The designer's RTL compiles successfully without errors.

Exit Criteria

The objective, mathematically provable goals that must be hit to sign off on verification:

- 100% of all P0 (critical) and P1 (major) test cases pass.
- 100% Code Coverage is achieved (Statement, Branch, Toggle, and FSM).
- 100% Functional Coverage is closed (as defined in the covergroups of the VPlan).
- Zero open P0/P1 bugs remain in the issue tracking system.
- No unexplained simulation failures or regression drops are observed over a 72-hour window.

9.6 Traceability

Traceability is the process of linking every requirement, test case, and coverage metric together to ensure there are no verification blind spots.



The Hazards of Poor Traceability

- **A test with no coverage item:** You might have written a test case to cover an edge case, but without a dedicated functional `coverpoint`, you have no mathematical proof that the simulator actually executed that path under physical conditions.
- **A coverage item with no test:** You have a covergroup in your verification code, but no active test scenario triggers it. It remains permanently empty, blocking 100% coverage closure.

The Traceability Matrix

The ultimate deliverable of a professional verification engineer is the **Traceability Matrix**. It compiles all architectural features, design specifications, tests, and coverage checkpoints into a single unified table:


RTL Spec Ref	Feature Description	Req-ID	Target Test Case	Functional Coverage Link	Status
Section 6.2	FIFO must assert <code>empty</code> when counter matches 0.	RST_001	tb_fifo_init.sv	cg_fifo.cp_empty_flag	Pass
Section 6.3	FIFO must block writes when <code>full</code> is active.	NEG_001	tb_fifo_overflow.sv	cg_fifo.cp_write_when_full	Pass
Section 6.4	Simultaneous read + write when FIFO is full.	CORNER_004	tb_fifo_simul_full.sv	cg_fifo.cross_full_rw	In Progress
Section 6.4	Simultaneous read + write when FIFO is empty.	CORNER_005	tb_fifo_simul_empty.sv	cg_fifo.cross_empty_rw	Not Started
Section 6.2	FIFO must assert <code>full</code> when counter matches <code>DEPTH</code> .	FUNC_002	tb_fifo_fill.sv	cg_fifo.cp_full_flag	Pass


Module 10: Verification Methodologies & Stimulus Strategies

This module focuses on the practical implementation of verification environments, defining checking mechanisms, selecting stimulus strategies, and prioritizing test plans for efficient coverage closure.

10.1 Checker Types

In hardware verification, a **checker** is the component responsible for determining whether the Device Under Test (DUT) is behaving correctly. Choosing the right checker type depends on the nature of the feature you are verifying.





1. Assertions (SystemVerilog Assertions - SVA)

An assertion is a mathematical property that must hold true at specific clock cycles or over defined temporal windows during simulation.

- **Best for:** Verifying protocol invariants, interface timing rules, and critical signal relationships that must *always* hold true.
- **Why it shines:** They run passively and trigger automatically. If a violation occurs, the simulator flags the exact simulation cycle and the line of code that failed.
- **Examples:**
 - The `full` and `empty` flags must never be asserted high simultaneously.
 - The `seq_detected` output must stay low if the input sequence has not completed.
- **SVA Concept Syntax:** Code snippet

```
// Continuous safety check: full and empty can never be high together
assert property (@(posedge clk) !(full && empty));
```

2. Scoreboards

A scoreboard is a dynamic class or component that predicts the expected behavior of the DUT, monitors actual outputs, and performs automatic data integrity comparisons.

- **Best for:** Tracking data path correctness, data value integrity, and transactional order preservation (e.g., FIFOs, processors, or communication pipelines).
- **Golden Reference Model:** The scoreboard contains a simplified, high-level software representation of the hardware (e.g., a pure C++ or SystemVerilog queue class) to predict what the DUT *should* produce.
- **Example:** A FIFO scoreboard uses a software queue (`logic [7:0] expected_queue[$]`). When a write transaction is detected by the input monitor, the data is pushed onto the queue. When a read is detected by the output monitor, the scoreboard pops the oldest entry and verifies it matches the literal value on the DUT's `dout` port.

3. Checkers with Expected Values

A streamlined checker designed specifically for highly controlled, non-random environments.

- **Best for:** Directed corner-case test scenarios where the stimulus is hand-coded, and you know exactly what the output must be at a precise clock cycle.
- **Why it shines:** Simple to write. It does not require a complex dynamic reference model.
- **Example:** In a sequence detector testbench searching for `1011`, the testbench drives exactly `1`, `0`, `1`, `1` over four clock edges and asserts an immediate check.
- Code snippet:

```
@(posedge clk);
assert(seq_detected == 1'b1) else $error("Pattern 1011 failed to trigger match.");
```

10.2 Directed vs. Random Stimulus

The stimulus generator is the engine of your testbench. It can be constructed using either hand-coded steps or constrained mathematical randomization.

Directed Stimulus

In a directed test, the engineer manually writes the exact sequence of inputs cycle-by-cycle.

- **Best for:** Initial smoke tests, power-up/reset sequences, strict protocol handshakes, and targeted negative error testing.
- **Pros:** Highly predictable, easy to debug, and fast to code for simple operations.
- **Cons:** Extremely limited coverage breadth. It can only find bugs in the scenarios the engineer explicitly thought to write down.

Random / Constrained-Random Stimulus

The generator randomly selects input signals, data payloads, and arrival delays, constrained only by the basic legal limits of the interface protocol.

- **Best for:** Broad functional coverage, discovering highly complex multi-system interactions, and long-term stress testing.
- **Pros:** Discovers hidden "corner cases of corner cases"—complex state overlaps that a human designer or verifier would never think of.
- **Cons:** Harder to debug because a failure might occur after millions of randomized cycles. It strictly requires a robust scoreboard to dynamically evaluate the random outputs.
- **Constraint Example:** Generating random read and write requests while guaranteeing that writes are suppressed when the FIFO is full, and reads are suppressed when it is empty.
- Code snippet:

```
constraint c_fifo_ops {
    if (fifo_full) wr_en == 0;
    if (fifo_empty) rd_en == 0;
}
```

Choosing the Right Strategy

Scenario / Feature	Preferred Method	Reasoning
Power-On Reset	Directed	Predictable, simple, and required as an entry gate for all other tests.
Negative Protocol Testing	Directed	Forces illegal behavior precisely to verify error handling.
Data Path Throughput	Constrained-Random	Discovers bottlenecks and buffer overflows by varying packet gaps.
Protocol Invariant Rules	Assertion-Based	Active across both directed and random suites to detect interface violations.

10.3 Priority Levels in Test Plans

To manage schedule constraints and tape-out risk, verification requirements are categorized into strict priority tiers.

Priority 0 (P0) — Must Pass (Critical Path)

Without these features passing, the design is fundamentally broken and cannot be signed off. If a P0 test fails, tape-out is instantly blocked.

- **Examples:** Core chip reset operates cleanly, basic register read/write works, sequential pattern detection matches, and FIFO full/empty flags assert correctly.

Priority 1 (P1) — High Importance

Features that represent the standard operational capabilities of the chip. A failing P1 test blocks sign-off, but under extreme schedule pressure, workarounds (like software-level bypasses) might be evaluated.

- **Examples:** Simultaneous read and write operations at boundary limits, or sustained clock cycles of back-to-back operations without data corruption.

Priority 2 (P2) — Nice to Have

Unusual, highly specific, or extreme operational scenarios. Passing these adds massive confidence but does not threaten tape-out schedules.

- **Examples:** Multi-day continuous stress tests, extremely rare input noise patterns, or complex parameter overrides (e.g., simulating a 1-bit wide data bus).

10.4 Matching Features to Methods

The core of a professional verification plan is matching each functional feature with its most efficient check and stimulus method:

Feature Type	Stimulus Method	Checker Method
Protocol Invariants (Static)	Any (All Tests)	Assertions (SVA)
Data Correctness (Dynamic)	Constrained-Random	Scoreboard
Specific Handshake Sequence	Directed	Checker with Expected Value
Broad Functional Reach	Constrained-Random	Scoreboard + Covergroups
Negative Error Handling	Directed	Assertions (Verify error fires)

Module 11: Functional & Code Coverage

Coverage provides an objective, mathematical answer to the fundamental verification question: **"How do we know when we are done testing?"** It converts subjective confidence into concrete metrics required for silicon sign-off.

11.1 Why Coverage Exists

Without coverage metrics, testing ends when the project schedule runs out of time—an unacceptably high risk for tape-out.

- **The Baseline Metric:** High coverage does not guarantee a bug-free design, but **low coverage guarantees insufficient testing**.
- **Objective Sign-off:** It provides verifiable proof to project leads and sign-off reviewers that every critical operational state, scenario, and boundary condition has been thoroughly exercised.

11.2 Code Coverage (Implicit Metrics)

Automatically collected by the simulation tool, code coverage tracks whether the physical RTL code structures were executed during simulation.

- **Statement Coverage:** Verifies if every line of RTL executed at least once.
- **Branch Coverage:** Checks if every path of conditional logic (`if-else` , `case`) was traversed.
- **Toggle Coverage:** Ensures every bit of a net or register flipped from \$0 \rightarrow 1\$ and \$1 \rightarrow 0\$.
- **FSM Coverage:**
 - *State Coverage:* Evaluates if every defined state was reached.
 - *Transition Coverage:* Evaluates if every legal state-to-state arc was traversed (critical for uncovering state machine bugs).
- **Condition Coverage:** Analyzes compound boolean expressions (e.g., `A && B`) to confirm each sub-condition evaluated to true and false independently.

Target: 100% Code Coverage is the minimum baseline for sign-off. Anything less points to unverified paths or dead RTL code.

11.3 Functional Coverage (Explicit Metrics)

Written explicitly by the DV engineer using SystemVerilog constructs (`covergroup`), functional coverage tracks high-level design intent and scenarios that the simulator cannot infer from code alone.

Code snippet:


```
covergroup cg_fifo @(posedge clk);
  // 1. Coverpoint: Track specific buffer counts
  cp_count: coverpoint count {
    bins empty = {0};
    bins full  = {DEPTH};
    bins mid   = {[1:DEPTH-1]};
  }

  // 2. Cross Coverage: Check simultaneous operational states
  cross_rw: cross wr_en, rd_en;
endgroup
```

- **Coverpoint:** Tracks a specific variable, signal, or state expression.
- **Bins:** Defined value sets or range buckets within a coverpoint that hold explicit operational meaning.
- **Cross Coverage:** Tracks combinations of two or more coverpoints occurring concurrently (e.g., simultaneous write and read assertions).
- **Assertion Coverage:** Measures how many times a given assertion's temporal property evaluated successfully (`cover property`).

11.4 Coverage Closure

Coverage closure is the systematic engineering process of driving all coverage metrics to target goals (typically **90–100%**).



1. Run constrained-random test suites across different seed variations.
2. Generate simulator coverage reports to locate **Coverage Holes** (unhit bins).
3. Write targeted, **directed tests** specifically engineered to force the DUT into those unhit boundary conditions.

11.5 Linking Coverage to Tests & Traceability

Every covergroup item should explicitly map back to a specific **Req-ID** in the Verification Plan.

- **Closed Traceability Loop:**

\$\$\text{RTL Feature} \rightarrow \text{Req-ID (e.g., FUNC_003)} \rightarrow \text{Test Case} \rightarrow \text{Coverage Link}\$\$\$

- A coverage item without a linked test means a scenario was never deliberately targeted—it relies entirely on accidental random hits.
- Including explicit **Coverage Links** in the VPlan ensures total accountability during sign-off audits.

11.6 Designing a Coverage Model

Build your coverage model methodically by referencing the initial feature list:


Design Type	Coverage Strategy	Implementation Focus
Control / FSM	State & Transition Bins	Create explicit bins for every legal state and arc transition.
Datapath	Boundary & Range Bins	Target extreme values: <code>0</code> , <code>1</code> , <code>MAX-1</code> , <code>MAX</code> , and mid-range partitions.
Data Buses	Pattern Bins	Track toggles for all-zeros (<code>'0</code>), all-ones (<code>'1</code>), and alternating bit patterns (<code>'hA</code> , <code>'h5</code>).
Interfaces	Cross Bins	Cross concurrent control signals (e.g., <code>valid \$ \times \$ ready</code> , <code>wr_en \$ \times \$ rd_en</code>).

Module 12: Traceability Audit & Sign-Off Closure Metrics

This module focuses on executing a complete **Traceability Check** to ensure an unbroken verification chain and adopting a **Closure Mindset** to prepare a design for final verification execution.

12.1 The Unbroken Traceability Chain

Before executing a single simulation or declaring a verification plan complete, you must conduct a bi-directional audit of the entire verification infrastructure. This process verifies that every functional requirement maps directly to measurable verification components.



The 4-Step Traceability Audit Checklist

- Spec Feature $\rightarrow$ Req-ID:** Read the RTL specification line-by-line (Interface Table, FSM Transitions, Corner Cases). Confirm that every functional capability maps to at least one unique **Req-ID** in the feature matrix.
- Req-ID $\rightarrow$ Coverage Item:** Verify that every **Req-ID** specifies a concrete checker or functional coverage target (e.g., a specific SVA property, dynamic scoreboard check, or SystemVerilog `coverpoint` bin).
- Coverage Item $\rightarrow$ Req-ID Link:** Inspect the coverage plan in reverse. Every defined covergroup bin or cross-coverage item must cite the exact **Req-ID** responsible for exercising it.
- Identify Broken Links:**
 - Uncovered Feature:** A feature in the spec with no Req-ID. *Risk: Unverified hardware behavior entering production.*
 - Untracked Test:** A test case or Req-ID with no coverage item. *Risk: Simulation runs without objective proof that the scenario actually occurred.*
 - Orphaned Coverage:** A `coverpoint` with no corresponding Req-ID. *Risk: Wasted verification effort targeting an undefined requirement.*

12.2 Verification Metrics & Sign-Off Closure Thinking

Adopting a "closure mindset" means evaluating your verification plan using objective sign-off criteria rather than subjective completion estimates.

Priority Level Audit (P0 / P1 / P2)


- **P0 Audit (Tape-Out Blockers):** Ensure 100% of P0 requirements (power-on reset, fundamental read/write paths, basic state transitions) have explicit test cases, direct assertions, and high-priority coverage models assigned.
- **P1 & P2 Alignment:** Verify that complex interaction tests (e.g., simultaneous read and write actions at full or empty boundaries) are properly tagged as P1, while extreme parameter overrides or long stress runs are placed under P2.

Corner-Case Cross-Check

Revisit the **Corner Case Registry** from your RTL Spec. Every single edge scenario documented during the design phase—such as reset mid-transaction, illegal operational requests, or pointer wraparounds—must be explicitly represented as a dedicated Req-ID (typically under the `CORNER_` or `NEG_` categories).

12.3 Preparing for Testbench Implementation

Once the Traceability Audit is complete and all links are verified, the document handoff occurs. A DV engineer takes this completed Verification Plan and executes the following setup steps:



1. **Testbench Architecture Construction:** Build the verification environment framework (Interfaces, Clock/Reset Generators, Drivers, Monitors, and Scoreboards) to match the DUT ports listed in the Interface Table.
2. **Coverage Model Coding:** Implement SystemVerilog `covergroup` and `coverpoint` structures matching the explicit names defined in the VPlan.
3. **Test Case Authoring:** Write test scenarios targeting each specific Req-ID (starting with P0 directed reset/smoke tests, followed by constrained-random functional and stress suites).
4. **Regression & Coverage Closure Execution:** Run automated simulation runs, parse coverage reports, and iterate on directed tests to close any remaining coverage holes.